


Intro to Compute Canada

Evgeny Naumov



compute | **calcul**
canada | canada


Calcul Québec

Financial Partners



What is Compute Canada?

Overview

- Provides advanced research computing (ARC) resources to Canadian researchers
- Provides **support** to researchers using its infrastructure

Overview

- Operates in close cooperation with regional partners:
 - ACENET
 - Calcul Québec
 - Compute Ontario
 - WestGrid
- Funded by provincial, federal and university money

Common Environment

- Starting in 2017, there has been a push to standardize the software stack
- Goal is to present a common user interface across all assets
- This is the interface I will present, YMMV may vary on consortium servers

Bureaucracy

How to register?

- Register for an account (details by searching “compute canada account”)
 - At this point you can use limited RAS resources
- **Principal Investigators** (same eligibility as for federal grants) can get more resources for their groups through the RAC
 - Pls’ students can join their groups to access resources
 - Can be in more than one group - **choose the right account when running jobs!**

Glossary

- **RAS** - Rapid Access Service
 - 20% of compute; for small jobs; anyone with an account can use immediately
- **RRG** - Resources for Research Groups
 - 80% of compute; awarded on merit through RACs
- **RAC** - Resource Allocation Competition
 - Process by which RRG compute is allocated

Consortium Accounts

- To use consortium accounts (e.g. on `calculquebec.ca` servers), need to register separately
- “*Apply for Consortium Account*” in the Compute Canada account page
- Separate credentials
- Computing environments may differ a bit
 - In the process of standardization

The Basics

Logging in

- SSH into the appropriate server using either your
 - Compute Canada credentials (.computecanada.ca)
 - Consortium creds (e.g. .calculquebec.ca)
- SSH best practises apply:
 - Private keys; PK encryption
- **The login node is not a compute node**
 - If you run calculations on it you will be ~~federally prosecuted~~ gently reminded not to!

Storage

- There are three fundamental storage areas
 - `/home/yourusername` (`=: ~`)
 - On remote server, NFS = slow!
 - 10GB of space = not good for data!
 - Use for config files, documents, code, etc.
 - `~/scratch/`
 - Local to the server = fast!
 - 10 TB allocation
 - Use for... scratch data, rapid iteration
 - `/project/...`
 - RRG space, amount varies
 - Use for relevant project data

Running Jobs

- The login node is not a compute node!
- SLURM scheduler is used
- Several ways of using it
 - `sbatch` <your run script>
 - `srun` [flags] <your code>
 - `salloc` (for interactive session)
- Management commands
 - `squeue` check jobs
 - `scancel` cancel jobs
- Read the SLURM docs!

Running Jobs: Accounting

- Compute is metered!
- If you're in multiple projects:
 - Run the job under the **appropriate account**
 - `--account=` parameter to slurm
 - Your unix groups correspond to your accounts
 - Be careful with quota. Compute is limited, be efficient and considerate
- Conversely, don't use RAS (`def-xxx` account) if you have appropriate RRG (`rrg-xxx`). Slow!

There is much more to it

- Efficient use of HPC is hard
- Things to master:
 - When/how to use MPI
 - Efficient SLURM space and time requests (ask for enough, but no more)
 - `--mem`, `--n-tasks`, `--n-nodes`, `--time`, **etc.**
 - Checkpointing
 - Preemptible code
- It's worth taking some time to make sure you are not doing something inefficient!

GPU

- Not all servers have GPUs
 - `{cedar, graham}.compute.canada.ca`
 - `helios.calculquebec.ca`
 - *Others, consortium sites have lists*
- Use `--gres=gpu:1` (2, 4, ...) to request GPUs
- For distributed training:
 - Envvar **SLURM_JOB_NODELIST** gives node hostnames of allocated nodes
 - Use these to set up param servers, workers, etc.

Summary

- Register for account, join your PI's group
- Use SSH to log in
- Store code and config in `/home/yourname`
- Store big files in `~/scratch` **or** `/project/...`
- **NO BIG COMPUTE ON LOGIN NODES**
 - one-off tests/PoCs are fine
- Request GPU from SLURM with `--gres=gpu:1`

Software Installation

Primary Stacks

- **Lmod - module**

- `module` command
- Loads pre-configured software stacks
- Handles libs, compatibility automatically

- **Easybuild - eb**

- For building performance-sensitive scientific software not available through Lmod

- **Nix - nix-env**

- For “personal” software like editors, stream processors

- **Pip - pip**

- For python packages

Lmod

- The most important command
 - Selects compiler, interpreter, library versions
 - Automatically manages paths, etc
- Common commands:
 - `module load <name/version>` - load s/w
 - `module remove <name>`
 - `module av <name>` - quick search
 - `module spider <name>` - extensive search
- Implemented as bash functions, **if using other shell redefine them**
- Read the docs!

Easybuild

- An extensive, reproducible build system
- Usually used internally to create modules
- Can be used to build newer software versions
- Commands
 - `eb -S <name>` search for s/w
 - `eb <recipe>` build s/w
- Not needed often, but remember it's there!

Nix

- The package manager from nixos
- Rarely needed, nix packages already in default path
- Commands
 - `nix-env --install --attr <name>`
 - `nix-env --uninstall <name>`
 - `Nix-env --query --available <name>`
- CC has a fork of nixos nixpkgs:
 - <https://github.com/ComputeCanada/nixpkgs>
 - Feel free to clone and modify to your needs

Python

- **The golden rules:**

- **AVOID CONDA**

- Supported on a best-effort basis, discouraged
- It infests your .bashrc leading to lots of stupid [PYTHON]PATH related tickets

- **USE VIRTUALENVS**

```
$ module load python/x.y.x
```

```
$ mkdir venv
```

```
$ virtualenv venv
```

```
$ source venv/bin/activate
```

```
(venv) $ pip install foo
```


Python: wheels

- Wheels are precompiled python packages
- We have custom wheels for big packages
 - Compiled with fast libs and optimizations
 - *I'm working on a better build system to improve interoperability and keep it more up-to-date*
- Some important wheels:
 - `pip install torch torch_{cpu,gpu}==0.4.0`
 - `pip install tensorflow tensorflow-{cpu,gpu}`

Local Libs: `setrpaths.sh`

- Some software (e.g. MuJoCo) requires libraries to be installed locally.
- These have hardcoded **RPATHs** (shared object lookup paths), will die on CC environment
- The fix:
 - `setrpaths.sh --path <lib dir> [--add-origin]`
 - More info at \$(search for “installing software local directory” on wiki)

Summary

1. Check modules (`module {av, spider}`)
2. Check EasyBuild (`eb -S/eb`)
3. Python:
 - a. Conda = BAD
 - b. Virtualenv = GOOD
4. Nix for minor software (your favourite editor or stream processor, etc.)

Support

How to get help?

- First, read the wiki! It's a fine document with concrete instructions:
 - <https://docs.compute canada.ca>
 - Consortia have own wikis (e.g. wiki.calcul quebec.ca)
- Email support@compute canada.ca
 - That's me!
 - Tips:
 - Use Compute Canada account email address
 - Be specific
 - Give consent to look at your code (acct. option)

Things we resolve

- Admin/login/space/compute quota issues
- If pip-installed software crashes
 - Happens sometimes since our environment is nontrivial and we mix compilers
- If you need new/upgraded software
 - check easybuild first
- Access to ~~evil~~ proprietary software
- You're not sure how to implement something
 - Yes, we help you code your algos and run them on HPC
- Anything else you're stuck on!

Thank you and happy computing!